

Prompt Caching

Prompt Caching is a feature that let us save and reuse parts of your prompt across multiple API calls, so you don't resend the same content repeatedly.

Problem:

```
[System prompt - 2000 tokens] ← same every time
[Few-shot examples - 3000 tokens] ← same every time
[Documents/context - 5000 tokens] ← same every time
[User message - 50 tokens] ← changes each time
```

So without the prompt caching, the model has to process 10000+ tokens every single call, although some of them never changed

Example:

```
# Anthropic example
messages = [
    {
        "role": "user",
        "content": [
            {
                "type": "text",
                "text": "Here is a 10,000 word legal document...",
                "cache_control": {"type": "ephemeral"} # ← mark for caching
            },
            {
                "type": "text",
                "text": "Summarize section 3." # ← changes each call
            }
        ]
    }
]
```